

Lecture 25 EEP118

Limited Dependent Variable

1. Logit Model Parameters/ Marginal Effects

Parameters

Marginal effects for continuous x

Marginal effects for discrete x

2. Estimation Maximum Likelihood

3. Tests. Goodness of Fit. Likelihood Ratio Test

The chi square distribution

Guest speaker: Law school

Study all of chapter 17.1

Posted all remaining DA and solutions, Practice final also

Limited Dependent Variable Y

The basic context of this set of lectures is when Y is not continuous

$Y=0$ or 1 , Y is binary. YES/NO

Use a Data set on Women labor force participation

Source: MROZ.RAW in Wooldridge. T.A. Mroz (1987), "The Sensitivity of an Empirical Model of Married Women's Hours of Work to Economic and Statistical Assumptions," *Econometrica* 55, 765-799.

$Y=1$ or 0 column called `inlf` (short for in labor force)

Obs: $N=753$

`inlf` byte %9.0g `inlf=1` if in labor force, 1975, `inlf=0` otherwise

`age` byte %9.0g woman's age in years

`educ` byte %9.0g years of schooling

`kidslt6` byte %9.0g # kids < 6 years

kidsge6 byte %9.0g # kids 6-18

nwifeinc float %9.0g (faminc - wage*hours)/1000

hushrs int %9.0g hours worked by husband, 1975

husage byte %9.0g husband's age

huseduc byte %9.0g husband's years of schooling

huswage float %9.0g husband's hourly wage, 1975

city byte %9.0g =1 if live in SMSA

```
In [1]: # Load the 'pacman' package
library(pacman)
#packages to use load them now using the pacman "manager"
p_load(dplyr, haven, readr)
#Another great feature of p_load(): if you try to load a package that is not
p_load(ggplot2)

pacman::p_load(lfe, lmtest, haven, sandwich, tidyverse)
# lfe for running fixed effects regression
# lmtest for displaying robust SE in output table
# haven for loading in dta files
# sandwich for producing robust Var-Cov matrix
# tidyverse for manipulating data and producing plots

#The big difference with Stata that appears here is lm() by default
#doesn't compute robust SE - we have to use additional packages/functions
#to compute it. felm does allow for multi-way clustering by default though
#which is nice.

#I added an alternate version of the first plots to show that we can
#change the color of the points according to whether the prediction
#is in [0,1] or outside of it. You can also specify factor(inlf) for
#the latter plots of actual vs. predicted to only have the values 0 or 1 on
#the x-axis.

pacman::p_load(lfe, lmtest, margins, haven, sandwich, tidyverse)
# lfe for running fixed effects regression
# lmtest for displaying robust SE in output table
# haven for loading in dta files
# sandwich for producing robust Var-Cov matrix
# tidyverse for manipulating data and producing plots
install.packages(sandwich)
install.packages(lfe)
install.packages(lmtest)
install.packages(tidyverse)
library(sandwich)
library(lfe)
```

```

library(lmtest)
library(tidyverse)

#install for margins
#to get marginal effects from Logit
install.packages("mfx")
library(mfx)

# alternate plot theme for ggplot
theme_ed <- theme(
  legend.position = "bottom",
  panel.background = element_rect(fill = NA),
  # panel.border = element_rect(fill = NA, color = "grey75"),
  axis.ticks = element_line(color = "grey95", size = 0.3),
  panel.grid.major = element_line(color = "grey95", size = 0.3),
  panel.grid.minor = element_line(color = "grey95", size = 0.3),
  legend.key = element_blank())

```

Installing package into ‘/srv/r’
(as ‘lib’ is unspecified)

Warning message:

“package ‘margins’ is not available for this version of R

A version of this package for your version of R might be available elsewhere,

see the ideas at

<https://cran.r-project.org/doc/manuals/r-patched/R-admin.html#Installing-packages>”

Warning message in p_install(package, character.only = TRUE, ...):

“”

Warning message in library(package, lib.loc = lib.loc, character.only = TRUE, logical.return = TRUE, :

“there is no package called ‘margins’”

Warning message in pacman::p_load(lfe, lmtest, margins, haven, sandwich, tidyverse):

“Failed to install/load:

margins”

Installing package into ‘/srv/r’

(as ‘lib’ is unspecified)

Error in as.character(x): cannot coerce type 'closure' to vector of type 'character'

Traceback:

1. install.packages(sandwich)
2. grepl("[.]tar[.](gz|bz2|xz)\$", pkgs)

In [2]: *#load data*

```

mydata<- read_dta("Lecture24MR0Z.DTA")
#Summary stats inlf age educ kidslt6 kidsge6 nwifeinc hushrs husage huseduc
summary(mydata)

```

inlf	hours	kidslt6	kidsge6
Min. :0.0000	Min. : 0.0	Min. :0.0000	Min. :0.000
1st Qu.:0.0000	1st Qu.: 0.0	1st Qu.:0.0000	1st Qu.:0.000
Median :1.0000	Median : 288.0	Median :0.0000	Median :1.000
Mean :0.5684	Mean : 740.6	Mean :0.2377	Mean :1.353
3rd Qu.:1.0000	3rd Qu.:1516.0	3rd Qu.:0.0000	3rd Qu.:2.000
Max. :1.0000	Max. :4950.0	Max. :3.0000	Max. :8.000

age	educ	wage	repwage
Min. :30.00	Min. : 5.00	Min. : 0.1282	Min. :0.00
1st Qu.:36.00	1st Qu.:12.00	1st Qu.: 2.2626	1st Qu.:0.00
Median :43.00	Median :12.00	Median : 3.4819	Median :0.00
Mean :42.54	Mean :12.29	Mean : 4.1777	Mean :1.85
3rd Qu.:49.00	3rd Qu.:13.00	3rd Qu.: 4.9708	3rd Qu.:3.58
Max. :60.00	Max. :17.00	Max. :25.0000	Max. :9.98

NA's :325

hushrs	husage	huseduc	huswage
Min. : 175	Min. :30.00	Min. : 3.00	Min. : 0.4121
1st Qu.:1928	1st Qu.:38.00	1st Qu.:11.00	1st Qu.: 4.7883
Median :2164	Median :46.00	Median :12.00	Median : 6.9758
Mean :2267	Mean :45.12	Mean :12.49	Mean : 7.4822
3rd Qu.:2553	3rd Qu.:52.00	3rd Qu.:15.00	3rd Qu.: 9.1667
Max. :5010	Max. :60.00	Max. :17.00	Max. :40.5090

faminc	mtr	motheduc	fatheduc
Min. : 1500	Min. :0.4415	Min. : 0.000	Min. : 0.000
1st Qu.:15428	1st Qu.:0.6215	1st Qu.: 7.000	1st Qu.: 7.000
Median :20880	Median :0.6915	Median :10.000	Median : 7.000
Mean :23081	Mean :0.6789	Mean : 9.251	Mean : 8.809
3rd Qu.:28200	3rd Qu.:0.7215	3rd Qu.:12.000	3rd Qu.:12.000
Max. :96000	Max. :0.9415	Max. :17.000	Max. :17.000

unem	city	exper	nwifeinc
Min. : 3.000	Min. :0.0000	Min. : 0.00	Min. : -0.02906
1st Qu.: 7.500	1st Qu.:0.0000	1st Qu.: 4.00	1st Qu.:13.02504
Median : 7.500	Median :1.0000	Median : 9.00	Median :17.70000
Mean : 8.624	Mean :0.6428	Mean :10.63	Mean :20.12896
3rd Qu.:11.000	3rd Qu.:1.0000	3rd Qu.:15.00	3rd Qu.:24.46600
Max. :14.000	Max. :1.0000	Max. :45.00	Max. :96.00000

lwage	expersq
Min. : -2.0542	Min. : 0
1st Qu.: 0.8165	1st Qu.: 16
Median : 1.2476	Median : 81
Mean : 1.1902	Mean : 178
3rd Qu.: 1.6036	3rd Qu.: 225
Max. : 3.2189	Max. :2025

NA's :325

Fixing Problem 2, make sure predictions are between 0 and 1

Solution Problem 1 -

use a functional for the probability as a function $G(\cdot)$ of the x s that stays between 0 and 1

e.g., the Logit Model!

the ratio of exponents in the logit below is always between 0 and 1

$$\text{Prob}[Y=1 \mid x] = G(\beta_0 + \beta_1 x_1 + \beta_2 x_2 + \dots + \beta_k x_k)$$

Get a G that stays between 0 and 1, and **the Logit is**

$$\text{Prob}[Y=1 \mid x] = G(\beta_0 + \dots + \beta_k x_k) = \frac{e^{\beta_0 + \beta_1 x_1 + \beta_2 x_2 + \dots + \beta_k x_k}}{1 + e^{\beta_0 + \beta_1 x_1 + \beta_2 x_2 + \dots + \beta_k x_k}}$$

This ratio of exponentials is always between 0 and 1 no matter the betas and x s

$$\text{Prob}[Y=1 \mid X] = \Lambda(\beta_0 + \beta_1 x_1 + \beta_2 x_2 + \dots + \beta_k x_k)$$

24

In [3]: `##### Fixing Problem 2 so that predicted Y hats are less than 1 and greater than 0
In R, use the glm(formula, data, family = binomial(link = "logit")) function

logit <- glm(inlf ~ nwifeinc + educ + age + kidslt6 + kidsge6, mydata, family = binomial(link = "logit"))
summary(logit)`

Call:

`glm(formula = inlf ~ nwifeinc + educ + age + kidslt6 + kidsge6, family = binomial(link = "logit"), data = mydata)`

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	0.722993	0.788698	0.917	0.359
nwifeinc	-0.034891	0.007884	-4.426	9.62e-06 ***
educ	0.257965	0.040744	6.331	2.43e-10 ***
age	-0.057553	0.012737	-4.519	6.23e-06 ***
kidslt6	-1.484437	0.198013	-7.497	6.55e-14 ***
kidsge6	-0.066363	0.067856	-0.978	0.328

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 1029.75 on 752 degrees of freedom
Residual deviance: 908.37 on 747 degrees of freedom
AIC: 920.37

Number of Fisher Scoring iterations: 4

Cannot easily interpret parameters here,

next class estimate implied marginal effects given the above estimated Logit parameters

Parameters not very meaningful here. (they enter two exponentials to get Phat)

What we want is if say education changes by one, how does the Prob(y=1) change?

Logit Model Marginal Effects

For a continuous variable x_1 education for example:

Given that $P(y = 1) = \frac{e^{\beta_0 + \beta_1 x_1 + \dots + \beta_k x_k}}{1 + e^{\beta_0 + \beta_1 x_1 + \dots + \beta_k x_k}} = \frac{e^z}{1 + e^z}$ where

$$z = \beta_0 + \beta_1 x_1 + \dots + \beta_k x_k$$

Then
$$\frac{\partial P(y=1)}{\partial x_1} = \frac{\partial P(y=1)}{\partial z} \frac{\partial z}{\partial x_1} = \frac{e^z}{(1+e^z)^2} \frac{\partial z}{\partial x_1}$$

$$\Leftrightarrow \frac{\partial P(y=1)}{\partial x_1} = \frac{e^z}{(1+e^z)^2} \beta_1$$

Logit Model Marginal Effects (ME)


For a continuous variable x_1 education for example:

For education, given the estimates above, if education changes by one, then the ME on the Prob(y=1) is given by

$$\frac{e^{\hat{\beta}_0 + \hat{\beta}_1 educ + \dots + \hat{\beta}_k x_k}}{(1 + e^{\hat{\beta}_0 + \hat{\beta}_1 educ + \dots + \hat{\beta}_k x_k})^2} \hat{\beta}_1 = \frac{e^{0.72544 + 0.2576 educ + \dots - 0.0351 nwifefinc}}{(1 + e^{0.72544 + 0.2576 educ + \dots - 0.0351 nwifefinc})^2} 0.2576$$

Where we substitute the estimated beta hats.

Note that the ME depends on the starting point of educ and also on all the other x's.



Logit Model Marginal Effects (ME)


For a continuous variable x_1 education for example:

$$\frac{e^{\hat{\beta}_0 + \hat{\beta}_1 educ + \dots + \hat{\beta}_k x_k}}{(1 + e^{\hat{\beta}_0 + \hat{\beta}_1 educ + \dots + \hat{\beta}_k x_k})^2} \hat{\beta}_1 = \frac{e^{0.72544 + 0.2576 educ + \dots - 0.0351 nwifeinc}}{(1 + e^{0.72544 + 0.2576 educ + \dots - 0.0351 nwifeinc})^2} 0.2576$$

Where we substitute the estimated beta hats.

Note that the ME depends on the starting point of educ

and also on all the other x's.



For a continuous variable x_1 education for example: How does one report the marginal effects (ME) then given that it depends on x s and starting point?

Report it for a fictitious person that would have all x 's at the average, that is, for $(educ) = 12.2$, $(kids) = 0.238$ etc etc, all average of all x 's, in this case, ME education is 0.0537, or 5.37 percentage points ---see next cell on how to get estimated ME

```
In [4]: #Marginal Effects (ME) at Mean X
logitmfx(inlf ~ nwifeinc + educ + age + kidslt6 + kidsge6, lec24df, atmean =
```

```
Error in logitmfx(inlf ~ nwifeinc + educ + age + kidslt6 + kidsge6, lec24df,
: could not find function "logitmfx"
Traceback:
```

```
In [5]: #compute the average ME, not at the mean of all X's as above

#Average Marginal Effects
logitmfx(inlf ~ nwifeinc + educ + age + kidslt6 + kidsge6, lec24df, atmean =
```

```
Error in logitmfx(inlf ~ nwifeinc + educ + age + kidslt6 + kidsge6, lec24df,
: could not find function "logitmfx"
Traceback:
```

```
In [6]: #generate predictions

mydata <- mutate(mydata, log_fit = logit$fitted.values) # add in the logit

#Reproduce figures for logit
# no need to use the second approach as we're always within [0,1] with logit
# set data and aesthetics (x and y vars here since the same for all elements
ggplot(mydata, aes(x = educ, y = log_fit)) +
  # First add points, color determined by whether in or out of [0,1]
```

```

geom_point() + # add points
# add horizontal lines, width slightly wider, making partially transparent
geom_hline(yintercept=0, size = 1.4, alpha = 0.35, color = "red") + # add
geom_hline(yintercept=1, size = 1.4, alpha = 0.35, color = "red") + # add
# generate labels
labs(title = "Predicted Probability of Labor Force Participation and Educat
      subtitle = "Logit Model",
      x = "Years of Schooling",
      y = "Fitted Value (Probability of in Labor Force Participation)")

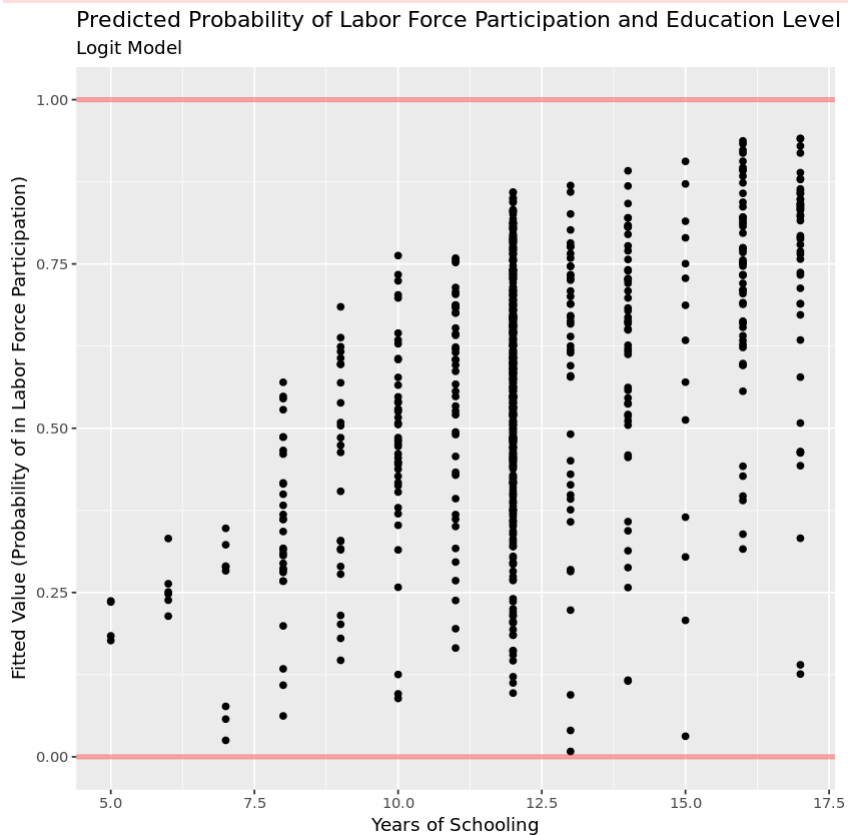
# actual vs predicted
ggplot(mydata, aes(x = factor(inlf), y = log_fit)) +
# First add points, color determined by whether in or out of [0,1]
geom_point() +
# add horizontal lines, width slightly wider, making partially transparent
geom_hline(yintercept=0, size = 1.4, alpha = 0.35, color = "red") + # add
geom_hline(yintercept=1, size = 1.4, alpha = 0.35, color = "red") + # add
# generate labels
labs(title = "Predicted vs Actual Probability of Labor Force Participation
      subtitle = "Logit Model",
      x = "Actual Labor Force Participation, 1975",
      y = "Estimated Labor Force Participation")

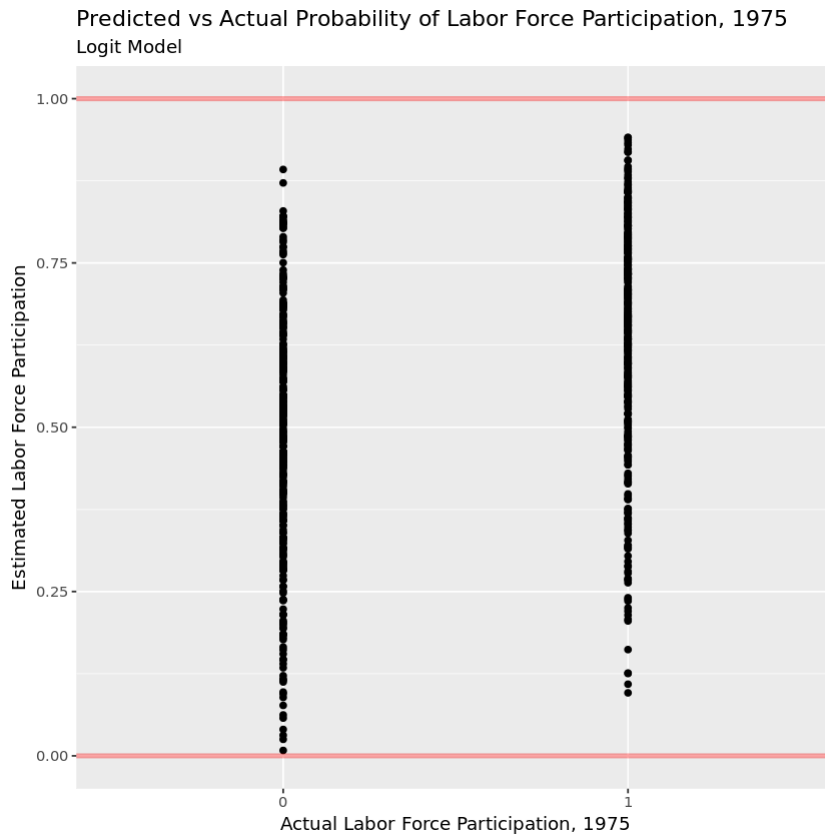
```

Warning message:

“Using `size` aesthetic for lines was deprecated in ggplot2 3.4.0.

i Please use `linewidth` instead.”





Marginal Effect for a Discrete variable X

For a discrete variable x_1 city for example:

We need to compute the difference in probability, that is $ME_{city} = \text{Prob}(y=1 | x, \text{city}=1) - \text{Prob}(y=1 | x, \text{city}=0)$

And once again we evaluate all at the average of all other x 's

(*) dy/dx is for discrete change of dummy variable from 0 to 1

```
In [7]: #run a logit with a city dummy variable
logit2 <- glm(inlf ~ nwifeinc + educ + age + kidslt6 + kidsge6+city, mydata,
summary(logit2)

#Average Marginal Effects
logitmfx(inlf ~ nwifeinc + educ + age + kidslt6 + kidsge6+city, lec24df, atr

#Marginal Effects at Mean X
logitmfx(inlf ~ nwifeinc + educ + age + kidslt6 + kidsge6+city, lec24df, atr
```

```
Call:
glm(formula = inlf ~ nwifeinc + educ + age + kidslt6 + kidsge6 +
    city, family = binomial(link = "logit"), data = mydata)
```

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	0.725440	0.789091	0.919	0.358
nwifeinc	-0.035075	0.008067	-4.348	1.37e-05 ***
educ	0.257560	0.040910	6.296	3.06e-10 ***
age	-0.057689	0.012800	-4.507	6.58e-06 ***
kidslt6	-1.484777	0.198075	-7.496	6.58e-14 ***
kidsge6	-0.066625	0.067901	-0.981	0.326
city	0.019103	0.174730	0.109	0.913

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 1029.75 on 752 degrees of freedom
 Residual deviance: 908.36 on 746 degrees of freedom
 AIC: 922.36

Number of Fisher Scoring iterations: 4

```
Error in logitmfx(inlf ~ nwifeinc + educ + age + kidslt6 + kidsge6 + city, :
could not find function "logitmfx"
Traceback:
```

For a city relative to not a city the probability of a woman being in the labor force increases by 0.004,

but not significantly because the p value of the marginal effect is 0.913

and confidence interval for city Marginal effect covers zero : lower= -0.06737453
 upper=0.075335006

Estimation of Logit - by Maximum Likelihood

Maximum Likelihood

Derivation: for each observation of a woman

Suppose woman i working $Y_i=1$, then, the prob is $\Pr(Y_i=1|x_i) = \Lambda(\beta_0 + \beta_1 x_{1i})$

Suppose woman j is not working, $Y_i=0$, then the prob of that is $\Pr(Y_j=0|x_j) = 1 - \Lambda(\beta_0 + \beta_1 x_{1j})$

Maximum Likelihood

The Probability of observing i working and j not is equal to the product below which is the

Likelihood

$$= (\Lambda \beta_0 + \beta_1 x_{1i}) * [1 - \Lambda(\beta_0 + \beta_1 x_{1j})]$$

$$= \Pr(Y_i=1|x_i) \text{ times } \Pr(Y_j=0|x_j)$$

- Put all the working in data together and all the non working
- The prob to see what we see in the sample is the product of the prob of all the working i's

$$\text{Likelihood} = \prod_i (\Lambda(\beta_0 + \beta_1 x_{1i})) \prod_j [1 - \Lambda(\beta_0 + \beta_1 x_{1j})]$$

$\text{all } Y_i = \text{inlf}_i = 1 \text{ if women in labor market}$

$Y_i = \text{inlf}_i = 0 \text{ if not in labor market}$

and the product of the prob of all the non working j's.

$$L = \prod_j \left[\frac{e^{X_j \beta}}{1 + e^{X_j \beta}} \right]^{y_j} \prod_i \left[1 - \frac{e^{X_i \beta}}{1 + e^{X_i \beta}} \right]^{1 - y_i}$$

21

- Logging all that

log Likelihood =

$$LL = \sum_i \ln[\Lambda(\beta_0 + X_i \beta)] + \sum_j \ln[1 - \Lambda(\beta_0 + X_j \beta)]$$

$\text{all } Y_i = \text{inlf}_i = 1 \text{ if women in labor market}$

$Y_i = \text{inlf}_i = 0 \text{ if not in labor market}$

$$\log L = \sum_j y_j * \log\left[\frac{e^{X_j \beta}}{1 + e^{X_j \beta}}\right] + \sum_i (1 - y_i) \log\left[1 - \frac{e^{X_i \beta}}{1 + e^{X_i \beta}}\right]$$

Estimation Logit, Max Likelihood

- Put all the working in data together and all the non working
- The prob to see what we see in the sample is the product of the prob of all the working i's and the product of the prob of all the non working j's.
- If we log all of that we get
- **log Likelihood =**

$$LL = \sum_i \ln[\Lambda(\beta_o + X_i \beta)] + \sum_j \ln[1 - \Lambda(\beta_o + X_j \beta)]$$

for all $Y_i = \text{inlf}_i = 1$ if women in labor market
for all $Y_i = \text{inlf}_i = 0$ if not in labor market

$$\log L = \sum_j y_i * \log\left[\frac{e^{X_i \beta}}{1 + e^{X_i \beta}}\right] + \sum_i (1 - y_i) \log\left[1 - \frac{e^{X_i \beta}}{1 + e^{X_i \beta}}\right]$$

In [8]: *#estimate a model with lots of X's*

```
logit_u <- glm(inlf ~ nwifeinc + educ + age + kidslt6 + kidsge6 + city + hushrs +
summary(logit_u)
```

Call:

```
glm(formula = inlf ~ nwifeinc + educ + age + kidslt6 + kidsge6 +
    city + hushrs + husage + huseduc + huswage, family = binomial(link = "logit"),
    data = mydata)
```

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)	
(Intercept)	2.1072318	0.9407917	2.240	0.0251	*
nwifeinc	-0.0182788	0.0128726	-1.420	0.1556	
educ	0.2893468	0.0478669	6.045	1.50e-09	***
age	-0.0383568	0.0224972	-1.705	0.0882	.
kidslt6	-1.5370349	0.2009480	-7.649	2.03e-14	***
kidsge6	-0.0648634	0.0684488	-0.948	0.3433	
city	0.0147352	0.1809473	0.081	0.9351	
hushrs	-0.0003818	0.0001706	-2.238	0.0252	*
husage	-0.0283468	0.0224390	-1.263	0.2065	
huseduc	-0.0354425	0.0365281	-0.970	0.3319	
huswage	-0.0434876	0.0372837	-1.166	0.2435	

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 1029.75 on 752 degrees of freedom
Residual deviance: 900.47 on 742 degrees of freedom
AIC: 922.47

Number of Fisher Scoring iterations: 4

What do you see in the output above?

AIC reported, a good measure of fit that is also used for model comparison

Akaike information Criterion (AIC) , not R squared any more, no more minimizing SSR

now we are maximizing log Likelihood as the estimation criterion, what are the parameters that make the sample we see the most likely?

AIC: 922.36

obtained by

Akaike Information Criterion

$AIC = \ln(ei^2/n) + (2k/n) = \ln(SSR/n) + (2k/n)$

Hypothesis testing for one coefficient?

```
In [9]: #Hypothesis testing for one coefficient
#Single parameter test- use normal z below

#Coefficients:
#
```

	Estimate	Std. Error	z value	Pr(> z)
\$(Intercept)	0.725440	0.789091	0.919	0.358
#nwifeinc	-0.035075	0.008067	-4.348	1.37e-05 ***
#educ	0.257560	0.040910	6.296	3.06e-10 ***
#age	-0.057689	0.012800	-4.507	6.58e-06 ***
#kidslt6	-1.484777	0.198075	-7.496	6.58e-14 ***
#kidsge6	-0.066625	0.067901	-0.981	0.326
#city	0.019103	0.174730	0.109	0.913

For example, reject that education coefficient is zero. z stat is 6.29 p value 3.06e-10 ***

Hypothesis Testing for multiple coefficients?

likelihood ratio test in step 2

and critical values of a chi squared distribution in step 3

Hypothesis Testing for multiple betas

LIKELIHOOD RATIO TEST

Example: $H_0: \beta_1 = \beta_2 = \beta_3 = \beta_4 = \dots = 0$ (q = number of restrictions)

Under the null hypothesis:

$LR = 2 [\text{Log likelihood unrestricted} - \text{Log Likelihood restricted}]$

is distributed $\chi^2(q)$ Chi square with q = degrees of freedom

The chi -square distribution and table




TABLE G.4

Critical Values of the Chi-Square Distribution

		Significance Level		
		.10	.05	.01
D e g r e e s o f	1	2.71	3.84	6.63
	2	4.61	5.99	9.21
	3	6.25	7.81	11.34
	4	7.78	9.49	13.28
	5	9.24	11.07	15.09
	6	10.64	12.59	16.81
	7	12.02	14.07	18.48
	8	13.36	15.51	20.09
	9	14.68	16.92	21.67
	10	15.99	18.31	23.21
F r e e d o m	11	17.28	19.68	24.72
	12	18.55	21.03	26.22
	13	19.81	22.36	27.69
	14	21.06	23.68	29.14
	15	22.31	25.00	30.58
	16	23.54	26.30	32.00
	17	24.77	27.59	33.41
	18	25.99	28.87	34.81
	19	27.20	30.14	36.19
	20	28.41	31.41	37.57
	21	29.62	32.67	38.93
	22	30.81	33.92	40.29
	23	32.01	35.17	41.64
	24	33.20	36.42	42.98
	25	34.38	37.65	44.31
	26	35.56	38.89	45.64
	27	36.74	40.11	46.96
	28	37.92	41.34	48.28
	29	39.09	42.56	49.59
	30	40.26	43.77	50.89

Example: The 5% critical value with $df = 8$ is 15.51.
Source: This table was generated using the Stata® function invchi.

5 Steps as usual in hypothesis Testing

STEPS in Hypothesis testing

- Specify the null and the alternative hypothesis
- Run logit with all x s on the right = unrestricted model
 - Get the Log Likelihood value for the unrestricted L_{UR}
- Then run logit omitting 4 x 's, we are testing whether those betas for those x 's are zero – this is the restricted model
 - Get the Log Likelihood value for the restricted L_R
- Compute Likelihood Ratio Test Statistic= $LR=2(L_{UR}-L_R)$
- Compare with critical value of χ^2 with 4 degrees of freedom for significance level chosen
- If critical value less than LR then we reject the null. Otherwise cannot reject the null

30

now coding and computing and doing the actual test

```
In [10]: #step 1 Null that coefficients on the four husbands characteristics, all four
#step 2

#likelihood testing

#run unrestricted model
logit_u <- glm(inlf ~ nwifeinc + educ + age + kidslt6 + kidsge6+city+hushrs+
summary(logit_u)

#get the log likelihood of the unrestricted model
```

```
Call:
glm(formula = inlf ~ nwifeinc + educ + age + kidslt6 + kidsge6 +
     city + hushrs + husage + huseduc + huswage, family = binomial(link = "logit"),
     data = mydata)
```

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)	
(Intercept)	2.1072318	0.9407917	2.240	0.0251	*
nwifeinc	-0.0182788	0.0128726	-1.420	0.1556	
educ	0.2893468	0.0478669	6.045	1.50e-09	***
age	-0.0383568	0.0224972	-1.705	0.0882	.
kidslt6	-1.5370349	0.2009480	-7.649	2.03e-14	***
kidsge6	-0.0648634	0.0684488	-0.948	0.3433	
city	0.0147352	0.1809473	0.081	0.9351	
hushrs	-0.0003818	0.0001706	-2.238	0.0252	*
husage	-0.0283468	0.0224390	-1.263	0.2065	
huseduc	-0.0354425	0.0365281	-0.970	0.3319	
huswage	-0.0434876	0.0372837	-1.166	0.2435	

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 1029.75 on 752 degrees of freedom
 Residual deviance: 900.47 on 742 degrees of freedom
 AIC: 922.47

Number of Fisher Scoring iterations: 4

```
In [11]: #run the restricted model
         #no husband charct as regressors

logit_r <- glm(inlf ~ nwifeinc + educ + age + kidslt6 + kidsge6+city, mydata)
summary(logit_r)

#get the log likelihood of restricted model
```



```
Call:
glm(formula = inlf ~ nwifeinc + educ + age + kidslt6 + kidsge6 +
     city, family = binomial(link = "logit"), data = mydata)
```

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	0.725440	0.789091	0.919	0.358
nwifeinc	-0.035075	0.008067	-4.348	1.37e-05 ***
educ	0.257560	0.040910	6.296	3.06e-10 ***
age	-0.057689	0.012800	-4.507	6.58e-06 ***
kidslt6	-1.484777	0.198075	-7.496	6.58e-14 ***
kidsge6	-0.066625	0.067901	-0.981	0.326
city	0.019103	0.174730	0.109	0.913

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 1029.75 on 752 degrees of freedom
 Residual deviance: 908.36 on 746 degrees of freedom
 AIC: 922.36

Number of Fisher Scoring iterations: 4

In [12]: *#get both log likelihood values for the test statistics we will compute to e*

```
#get log likelihood value unrestricted
logLik(logit_u)
```

'log Lik.' -450.2368 (df=11)

In [13]: *#get log likelihood value restricted*

```
logLik(logit_r)
```

'log Lik.' -454.1793 (df=7)

compute the chi square stat

By hand, you will do this in Pset 5:

$$LR = 2 (\text{loglikelihood UR} - \text{loglikelihood R}) = 2 * (-450.237 - + 454.179) = 2 * 3.94$$

$$\text{So } LR = \text{chi2}(4) = 7.89$$

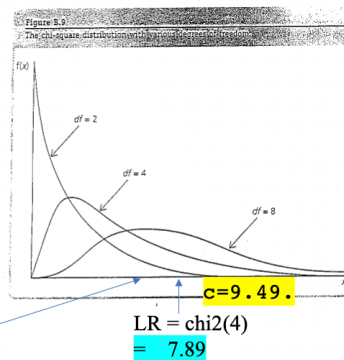
step 3 go to the table and get the critical value for a certain significance level

see below

Step 3: get critical value from Chi squared Table:

at 10% $c=7.78$

At 5% $c=9.49$



Step 4: get critical value from Chi squared Table: at 10% $c=7.78$. At 5% $c=9.49$. reject at 10% cannot reject at 5%.

Step 5: conclude with a sentence. At 5%, there is no statistical evidence that husbands characteristics matter for prob woman in labor force controlling for non wife income, kids, woman educ, etc

TABLE 8.4
Critical Values of the Chi-Square Distribution

	Significance Level		
	.10	.05	.01
1	2.71	3.84	6.63
2	4.61	5.99	9.21
3	6.25	7.81	11.34
4	7.78	9.49	13.28
5	9.24	11.07	15.09
6	10.64	12.59	16.81
7	12.02	14.07	18.48
8	13.36	15.51	20.09
9	14.68	16.92	21.67
10	15.99	18.31	23.21
11	17.28	19.68	24.72
12	18.55	21.03	26.22
13	19.81	22.36	27.69
14	21.06	23.68	29.14
15	22.31	25.00	30.58
16	23.54	26.30	32.00
17	24.77	27.59	33.41
18	25.99	28.87	34.81
19	27.20	30.14	36.19
20	28.41	31.41	37.57
21	29.62	32.67	38.93
22	30.81	33.92	40.29
23	32.01	35.17	41.64
24	33.20	36.42	42.98
25	34.38	37.65	44.31
26	35.56	38.89	45.64
27	36.74	40.11	46.96
28	37.92	41.34	48.28
29	39.09	42.56	49.59
30	40.29	43.77	50.89

Example: The 5% critical value with $df = 4$ is 9.49.
Source: This table was generated using the Stats® function levels.

step 4

at 10% $c=7.78 < LR=7.89$ so we reject the null at 10%

at 5% $c=9.49 > LR = 7.89$, so we cannot reject the null at 5%

Step5: conclude with a sentence. At 5%, there is no statistical evidence that husbands characteristics matter for prob woman in labor force controlling for non wife income, kids, woman educ, etc

all together

Step 1: $H_0 \text{ Beta_hushrs} = \text{Beta_husage} = \text{Beta_huseduc} = \text{Beta_huswage} = 0$

H_1 not H_0

Step 1: under the null 2 (loglikelihood UR - loglikelihood R) follows a Chi Square with q degrees of freedom

Step 2:

By hand, you will do this in Pset 5:

$LR = 2 (\text{loglikelihood UR} - \text{loglikelihood R}) = 2 * (-450.237 - + 454.179) = 2 * 3.94$

So $LR = \chi^2(4) = 7.89$

Step 3: get critical value from Chi squared Table: at 10% $c=7.78$ At 5% $c=9.49$. reject at 10% cannot reject at 5%.

Step 4/5: conclude with a sentence. At 5%, there is no statistical evidence that husbands charct matter for prob woman in labor force controlling for non wife income, kids, woman educ, etc

In [14]: *#in your career you can use a canned command, not in this class though...*
##in R: various equivalent specifications of the LR test
`lrtest(logit_u, logit_r)`

A anova: 2 × 5

	#Df	LogLik	Df	Chisq	Pr(>Chisq)
	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
1	11	-450.2368	NA	NA	NA
2	7	-454.1793	-4	7.885107	0.09587869

In R- for your future work in Metrics in life 😊

##in R: various equivalent specifications of the LR test

`lrtest(logit_u, logit_r)`

You get the output in R then:

Likelihood ratio test

Model 1: `inlf ~ nwifeinc + educ + age + kidslt6 + kidsge6 + city + hushrs + husage + huseduc + huswage`

Model 2: `inlf ~ nwifeinc + educ + age + kidslt6 + kidsge6 + city`

	#Df	LogLik	Df	Chisq	Pr(>Chisq)
1	11	-450.24			
2	7	-454.18	-3.94	7.8851	0.09588 .

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Here would not even reject at 10% because p value 0.0958

36

the end

In []: